

01_eda_proposal

December 5, 2025

1 Predicting Hospital Readmission Risk

1.1 Student A – Data Engineer / Preprocessing Lead (Binger)

1.2 1. Setup & Load

```
[2]: # 1.1 Imports & Display Settings
# -----
import os
import pandas as pd
import numpy as np

# Visualization
import missingno as msno
import matplotlib.pyplot as plt
import seaborn as sns
import re

# Plot style
plt.style.use("default")
sns.set()

# Pandas display options for easier inspection
pd.set_option("display.max_columns", 50)
pd.set_option("display.width", 120)

print("Libraries imported.")
```

Libraries imported.

```
[3]: !pip install missingno

/etc/zshenv:4: unmatched "
zsh:1: command not found: pip
```

```
[4]: from pathlib import Path

# Path to ../results/figures relative to the notebook
FIG_DIR = Path(..) / "results" / "figures"
```

```

FIG_DIR.mkdir(parents=True, exist_ok=True) # create if it doesn't exist

[5]: print("Dataset: Diabetes 130-US Hospitals for Years 1999-2008 (UCI Repository)")

Dataset: Diabetes 130-US Hospitals for Years 1999-2008 (UCI Repository)

[6]: # 1.2 Define Paths & Load the CSV
# -----
RAW_DATA_PATH = "../data/raw/diabetic_data.csv"

df = pd.read_csv(RAW_DATA_PATH)
print("Data loaded.")

```

Data loaded.

1.3 2. Review the Dataset

```

[7]: # 2.1 Quick Peek
# -----
# First 5 rows
df.head()

[7]:   encounter_id patient_nbr           race gender      age weight
admission_type_id discharge_disposition_id \
0            2278392        8222157    Caucasian Female  [0-10)    ?
6                  25
1            149190        55629189    Caucasian Female  [10-20)    ?
1                  1
2            64410         86047875 AfricanAmerican Female  [20-30)    ?
1                  1
3            500364        82442376    Caucasian   Male  [30-40)    ?
1                  1
4            16680         42519267    Caucasian   Male  [40-50)    ?
1                  1

           admission_source_id time_in_hospital payer_code      medical_specialty
num_lab_procedures num_procedures \
0                      1                 1      ?  Pediatrics-Endocrinology
41                     0
1                      7                 3      ?
59                     0
2                      7                 2      ?
11                     5
3                      7                 2      ?
44                     1
4                      7                 1      ?
51                     0

```

	num_medications	number_outpatient	number_emergency	number_inpatient	
diag_1	diag_2	diag_3	number_diagnoses	\	
0		1	0	0	0
250.83	?	?	1		
1		18	0	0	0
276	250.01	255	9		
2		13	2	0	1
648	250	V27	6		
3		16	0	0	0
8	250.43	403	7		
4		8	0	0	0
197	157	250	5		

	max_glu_serum	A1Cresult	metformin	repaglinide	nateglinide	chlorpropamide	glimepiride	acetohexamide	glipizide	\
0	NaN	NaN	No	No	No	No	No	No	No	
No	No	No								
1	NaN	NaN	No	No	No	No	No	No	No	
No	No	No								
2	NaN	NaN	No	No	No	No	No	No	No	
No	No	Steady								
3	NaN	NaN	No	No	No	No	No	No	No	
No	No	No								
4	NaN	NaN	No	No	No	No	No	No	No	
No	No	Steady								

	glyburide	tolbutamide	pioglitazone	rosiglitazone	acarbose	miglitol	troglitazone	tolazamide	examide	citoglipton	\
0	No	No	No	No	No	No	No	No	No	No	
No	No	No	No	No	No	No	No	No	No	No	
1	No	No	No	No	No	No	No	No	No	No	
No	No	No	No	No	No	No	No	No	No	No	
2	No	No	No	No	No	No	No	No	No	No	
No	No	No	No	No	No	No	No	No	No	No	
3	No	No	No	No	No	No	No	No	No	No	
No	No	No	No	No	No	No	No	No	No	No	
4	No	No	No	No	No	No	No	No	No	No	
No	No	No	No	No	No	No	No	No	No	No	

	insulin	glyburide-metformin	glipizide-metformin	glimepiride-pioglitazone	metformin-rosiglitazone	\
0	No	No	No	No	No	No
No						
1	Up		No	No	No	No
No						
2	No		No	No	No	No
No						

3	Up	No	No	No
No				
4	Steady	No	No	No
No				
	metformin-pioglitazone change	diabetesMed	readmitted	
0	No	No	No	NO
1	No	Ch	Yes	>30
2	No	No	Yes	NO
3	No	Ch	Yes	NO
4	No	Ch	Yes	NO

```
[8]: df.shape
df.dtypes
```

```
[8]: encounter_id           int64
patient_nbr            int64
race                  object
gender                 object
age                   object
weight                object
admission_type_id     int64
discharge_disposition_id int64
admission_source_id    int64
time_in_hospital      int64
payer_code              object
medical_specialty     object
num_lab_procedures    int64
num_procedures         int64
num_medications        int64
number_outpatient      int64
number_emergency       int64
number_inpatient       int64
diag_1                 object
diag_2                 object
diag_3                 object
number_diagnoses       int64
max_glu_serum          object
A1Cresult              object
metformin               object
repaglinide             object
nateglinide             object
chlorpropamide          object
glimepiride             object
acetohexamide           object
glipizide               object
glyburide               object
```

```

tolbutamide          object
pioglitazone        object
rosiglitazone        object
acarbose             object
miglitol              object
troglitazone         object
tolazamide            object
examide               object
citoglipton           object
insulin                object
glyburide-metformin   object
glipizide-metformin   object
glimepiride-pioglitazone object
metformin-rosiglitazone object
metformin-pioglitazone object
change                 object
diabetesMed            object
readmitted             object
dtype: object

```

```
[9]: # 2.2 Shape and Column List
print("Number of rows:", df.shape[0])
# -----
print("Number of columns:", df.shape[1])
print("\nColumns:")
print(df.columns.tolist())
```

Number of columns: 50

Columns:

```
['encounter_id', 'patient_nbr', 'race', 'gender', 'age', 'weight',
'admission_type_id', 'discharge_disposition_id', 'admission_source_id',
'time_in_hospital', 'payer_code', 'medical_specialty', 'num_lab_procedures',
'num_procedures', 'num_medications', 'number_outpatient', 'number_emergency',
'number_inpatient', 'diag_1', 'diag_2', 'diag_3', 'number_diagnoses',
'max_glu_serum', 'A1Cresult', 'metformin', 'repaglinide', 'nateglinide',
'chlorpropamide', 'glimepiride', 'acetohexamide', 'glipizide', 'glyburide',
'tolbutamide', 'pioglitazone', 'rosiglitazone', 'acarbose', 'miglitol',
'troglitazone', 'tolazamide', 'examide', 'citoglipton', 'insulin', 'glyburide-
metformin', 'glipizide-metformin', 'glimepiride-pioglitazone', 'metformin-
rosiglitazone', 'metformin-pioglitazone', 'change', 'diabetesMed', 'readmitted']
```

```
[10]: # 2.3 Data Types + Non-null Counts
# -----
# Quick infor summary: non-null counts and dtypes
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 101766 entries, 0 to 101765
```

Data columns (total 50 columns):

#	Column	Non-Null Count	Dtype
0	encounter_id	101766 non-null	int64
1	patient_nbr	101766 non-null	int64
2	race	101766 non-null	object
3	gender	101766 non-null	object
4	age	101766 non-null	object
5	weight	101766 non-null	object
6	admission_type_id	101766 non-null	int64
7	discharge_disposition_id	101766 non-null	int64
8	admission_source_id	101766 non-null	int64
9	time_in_hospital	101766 non-null	int64
10	payer_code	101766 non-null	object
11	medical_specialty	101766 non-null	object
12	num_lab_procedures	101766 non-null	int64
13	num_procedures	101766 non-null	int64
14	num_medications	101766 non-null	int64
15	number_outpatient	101766 non-null	int64
16	number_emergency	101766 non-null	int64
17	number_inpatient	101766 non-null	int64
18	diag_1	101766 non-null	object
19	diag_2	101766 non-null	object
20	diag_3	101766 non-null	object
21	number_diagnoses	101766 non-null	int64
22	max_glu_serum	5346 non-null	object
23	A1Cresult	17018 non-null	object
24	metformin	101766 non-null	object
25	repaglinide	101766 non-null	object
26	nateglinide	101766 non-null	object
27	chlorpropamide	101766 non-null	object
28	glimepiride	101766 non-null	object
29	acetohexamide	101766 non-null	object
30	glipizide	101766 non-null	object
31	glyburide	101766 non-null	object
32	tolbutamide	101766 non-null	object
33	pioglitazone	101766 non-null	object
34	rosiglitazone	101766 non-null	object
35	acarbose	101766 non-null	object
36	miglitol	101766 non-null	object
37	troglitazone	101766 non-null	object
38	tolazamide	101766 non-null	object
39	examide	101766 non-null	object
40	citoglipton	101766 non-null	object
41	insulin	101766 non-null	object
42	glyburide-metformin	101766 non-null	object
43	glipizide-metformin	101766 non-null	object
44	glimepiride-pioglitazone	101766 non-null	object

```

45 metformin-rosiglitazone    101766 non-null  object
46 metformin-pioglitazone    101766 non-null  object
47 change                      101766 non-null  object
48 diabetesMed                 101766 non-null  object
49 readmitted                  101766 non-null  object
dtypes: int64(13), object(37)
memory usage: 38.8+ MB

```

```
[11]: # 2.4 Basic Numeric Summary
# -----
df.describe()
```

```

[11]:      encounter_id  patient_nbr  admission_type_id  discharge_disposition_id
admission_source_id  time_in_hospital \
count   1.017660e+05  1.017660e+05      101766.000000          101766.000000
101766.000000      101766.000000
mean    1.652016e+08  5.433040e+07      2.024006          3.715642
5.754437          4.395987
std     1.026403e+08  3.869636e+07      1.445403          5.280166
4.064081          2.985108
min     1.252200e+04  1.350000e+02      1.000000          1.000000
1.000000          1.000000
25%    8.496119e+07  2.341322e+07      1.000000          1.000000
1.000000          2.000000
50%    1.523890e+08  4.550514e+07      1.000000          1.000000
7.000000          4.000000
75%    2.302709e+08  8.754595e+07      3.000000          4.000000
7.000000          6.000000
max    4.438672e+08  1.895026e+08      8.000000          28.000000
25.000000          14.000000

      num_lab_procedures  num_procedures  num_medications  number_outpatient
number_emergency  number_inpatient \
count           101766.000000  101766.000000      101766.000000          101766.000000
101766.000000      101766.000000
mean        43.095641       1.339730      16.021844          0.369357
0.197836          0.635566
std         19.674362       1.705807      8.127566          1.267265
0.930472          1.262863
min        1.000000       0.000000      1.000000          0.000000
0.000000          0.000000
25%        31.000000       0.000000      10.000000          0.000000
0.000000          0.000000
50%        44.000000       1.000000      15.000000          0.000000
0.000000          0.000000
75%        57.000000       2.000000      20.000000          0.000000
0.000000          1.000000

```

```

max           132.000000      6.000000      81.000000      42.000000
76.000000      21.000000

```

```

        number_diagnoses
count      101766.000000
mean       7.422607
std        1.933600
min        1.000000
25%        6.000000
50%        8.000000
75%        9.000000
max       16.000000

```

```
[12]: # Basic stats including categorical columns
df.describe(include="all").T.head(15) # first 15 rows of summary
```

			count	unique	top	freq	mean
std	min	25%	\				
encounter_id			101766.0	NaN	NaN	NaN	165201645.622978
102640295.983457	12522.0	84961194.0					
patient_nbr			101766.0	NaN	NaN	NaN	54330400.694947
38696359.346534	135.0	23413221.0					
race			101766	6	Caucasian	76099	NaN
NaN	NaN	NaN					
gender			101766	3	Female	54708	NaN
NaN	NaN	NaN					
age			101766	10	[70-80)	26068	NaN
NaN	NaN	NaN					
weight			101766	10	?	98569	NaN
NaN	NaN	NaN					
admission_type_id			101766.0	NaN	NaN	NaN	2.024006
1.445403	1.0	1.0					
discharge_disposition_id			101766.0	NaN	NaN	NaN	3.715642
5.280166	1.0	1.0					
admission_source_id			101766.0	NaN	NaN	NaN	5.754437
4.064081	1.0	1.0					
time_in_hospital			101766.0	NaN	NaN	NaN	4.395987
2.985108	1.0	2.0					
payer_code			101766	18	?	40256	NaN
NaN	NaN	NaN					
medical_specialty			101766	73	?	49949	NaN
NaN	NaN	NaN					
num_lab_procedures			101766.0	NaN	NaN	NaN	43.095641
19.674362	1.0	31.0					
num_procedures			101766.0	NaN	NaN	NaN	1.33973
1.705807	0.0	0.0					
num_medications			101766.0	NaN	NaN	NaN	16.021844

8.127566	1.0	10.0	
		50%	75%
encounter_id	152388987.0	230270887.5	443867222.0
patient_nbr	45505143.0	87545949.75	189502619.0
race	NaN	NaN	NaN
gender	NaN	NaN	NaN
age	NaN	NaN	NaN
weight	NaN	NaN	NaN
admission_type_id	1.0	3.0	8.0
discharge_disposition_id	1.0	4.0	28.0
admission_source_id	7.0	7.0	25.0
time_in_hospital	4.0	6.0	14.0
payer_code	NaN	NaN	NaN
medical_specialty	NaN	NaN	NaN
num_lab_procedures	44.0	57.0	132.0
num_procedures	1.0	2.0	6.0
num_medications	15.0	20.0	81.0

```
[13]: # 2.5 How Many Unique Values per Column?
# -----
# Number of unique values per column
nunique = df.nunique().sort_values(ascending=False)
nunique.head(20)
```

```
[13]: encounter_id          101766
patient_nbr            71518
diag_3                  790
diag_2                  749
diag_1                  717
num_lab_procedures      118
num_medications         75
medical_specialty       73
number_outpatient        39
number_emergency        33
discharge_disposition_id 26
number_inpatient         21
payer_code                18
admission_source_id      17
number_diagnoses        16
time_in_hospital         14
age                      10
weight                   10
admission_type_id        8
num_procedures           7
dtype: int64
```

```
[14]: # 2.6 Check the Target Variable (readmitted)
# -----
df["readmitted"].value_counts()
```

```
[14]: readmitted
NO      54864
>30     35545
<30     11357
Name: count, dtype: int64
```


```
[15]: # And with proportions
df["readmitted"].value_counts(normalize=True)
```

```
[15]: readmitted
NO      0.539119
>30     0.349282
<30     0.111599
Name: proportion, dtype: float64
```


```
[16]: # 2.7 Missing Values Overview
# -----
# How many "?" entries overall?
(df == "?").sum().sort_values(ascending=False).head(20)
```

```
[16]: weight              98569
medical_specialty       49949
payer_code              40256
race                   2273
diag_3                 1423
diag_2                  358
diag_1                  21
encounter_id            0
tolazamide               0
glyburide                0
tolbutamide               0
pioglitazone             0
rosiglitazone             0
acarbose                  0
miglitol                  0
troglitazone               0
citoglipiton               0
examide                  0
acetohexamide              0
insulin                  0
dtype: int64
```

```
[17]: msno.matrix(df)
plt.show()
```



```
[18]: msno.bar(df, sort='descending', color="#66a9bc")
plt.show()
```



Observations:

- The dataset contains 101,766 encounters and 50 attributes.
- The target variable `readmitted` is imbalanced, with the majority of encounters not readmitted within 30 days.
- Several columns (`weight`, `max_glu_serum`, `A1Cresult`, `medical_specialty`, `payer_code`) have a high fraction of missing values and will be removed in preprocessing.


```
[19]: # gender was coded differently so we use a custom count for this one
print('gender', df['gender'][df['gender'] == 'Unknown/Invalid'].count())
```

gender 3

```
[20]: # Drop invalid gender rows
df = df[df['gender'] != 'Unknown/Invalid']
print("Remaining rows:", len(df))
```

Remaining rows: 101763

```
[21]: # Bar plot of readmitted distribution
sns.countplot(x="readmitted", data=df)
plt.title("Distribution of Readmission Labels")
plt.xlabel("Readmitted category")
plt.ylabel("Count")
plt.tight_layout()
```



1.4 3. Data Cleaning

```
[22]: missing_ratio = df.isna().sum() / len(df)
print("Top columns with missing values:")
print(missing_ratio.sort_values(ascending=False).head(10))
```

```
Top columns with missing values:
max_glu_serum      0.947466
A1Cresult          0.832768
encounter_id        0.000000
nateglinide         0.000000
glimepiride         0.000000
acetohexamide       0.000000
glipizide           0.000000
glyburide           0.000000
tolbutamide          0.000000
pioglitazone        0.000000
dtype: float64
```

```
[23]: # Replace '?' with NaN for consistency
df = df.replace('?', np.nan)

# Drop columns with too many missing values or IDs
cols_to_drop = ['weight', 'payer_code', 'medical_specialty',
                 'encounter_id', 'patient_nbr']
df.drop(columns=cols_to_drop, inplace=True)

# Check duplicates
print("Duplicates:", df.duplicated().sum())

# Re-check missing data
df.isnull().sum().sort_values(ascending=False).head(10)
```

```
Duplicates: 0
```

```
[23]: max_glu_serum    96417
A1Cresult        84745
race              2271
diag_3            1423
diag_2            358
diag_1            21
tolbutamide       0
pioglitazone     0
rosiglitazone    0
acarbose          0
dtype: int64
```

Columns with excessive missing values (weight, payer_code, medical_specialty) or identifiers (encounter_id, patient_nbr) were removed to improve model stability.

```
[24]: # Check duplicates
print("Duplicates:", df.duplicated().sum())

# Re-check missing data
df.isnull().sum().sort_values(ascending=False).head(10)
```

Duplicates: 0

```
[24]: max_glu_serum    96417
A1Cresult        84745
race            2271
diag_3          1423
diag_2           358
diag_1            21
tolbutamide         0
pioglitazone         0
rosiglitazone         0
acarbose            0
dtype: int64
```

```
[25]: # Drop heavy missing-value columns
cols_to_drop = ['max_glu_serum', 'A1Cresult']
df.drop(columns=cols_to_drop, inplace=True)
print(f"Dropped columns: {cols_to_drop}")
```

Dropped columns: ['max_glu_serum', 'A1Cresult']

```
[26]: # Re-check missing data
df.isnull().sum().sort_values(ascending=False).head(10)
```

```
[26]: race            2271
diag_3          1423
diag_2           358
diag_1            21
tolazamide         0
tolbutamide         0
pioglitazone         0
rosiglitazone         0
acarbose            0
miglitol            0
dtype: int64
```

```
[27]: # Handle small missing subsets
# Fill missing 'race' with mode (most common value)
df['race'] = df['race'].fillna(df['race'].mode()[0])

# Fill diagnosis codes with a placeholder
for col in ['diag_1', 'diag_2', 'diag_3']:
```

```

df[col] = df[col].fillna('Unknown')

[28]: # Re-check missing data
df.isnull().sum().sort_values(ascending=False).head(10)

```

```

[28]: race          0
examide        0
glyburide      0
tolbutamide    0
pioglitazone   0
rosiglitazone  0
acarbose       0
miglitol        0
troglitazone   0
tolazamide      0
dtype: int64

```

Features `max_glu_serum` and `A1Cresult` were dropped due to excessive missingness ($>80\%$). Minor missing values in `race` were imputed with the most frequent category, while missing diagnosis codes were replaced with “Unknown” to retain encounters.

1.5 4. Feature Encoding & Target Transformation

In this section, we perform the final preprocessing and feature transformation steps to prepare the dataset for modeling. This includes encoding the target variable, transforming categorical features into numeric form, and creating new derived features.

```

[29]: # 4.1 Encode Target Variable (readmitted_binary)
# -----
# Convert readmission categories ("NO", ">30", "<30") into a binary target ↴variable.

df["readmitted"].value_counts()

readmit_map = {'<30': 1, '>30': 0, 'NO': 0}
df['readmitted_binary'] = df['readmitted'].map(readmit_map).astype('int8')

# Check distribution
print(df['readmitted_binary'].value_counts())
print(df['readmitted_binary'].value_counts(normalize=True))

# Note: The target is highly imbalanced (~89% no readmission vs. 11% readmitted ↴within 30 days)

```

```

readmitted_binary
0    90406
1    11357
Name: count, dtype: int64
readmitted_binary

```

```
0    0.888398
1    0.111602
Name: proportion, dtype: float64
```

The target variable `readmitted_binary` represents whether a patient was readmitted within 30 days (1) or not (0). The dataset exhibits a significant class imbalance, which will be addressed in the modeling phase.

```
[30]: # 4.2 Convert Age Ranges to Midpoints (age_mid)
# -----
# Convert categorical age intervals (e.g., "[40-50)") into numeric midpoints (e.
# ↪g., 45).

age_map = {
    '[0-10)': 5, '[10-20)': 15, '[20-30)': 25, '[30-40)': 35,
    '[40-50)': 45, '[50-60)': 55, '[60-70)': 65, '[70-80)': 75,
    '[80-90)': 85, '[90-100)': 95
}
df['age_mid'] = df['age'].map(age_map)
```

The new numeric feature `age_mid` preserves ordering information from the categorical `age` variable while simplifying numerical analysis and visualization.

```
[31]: # 4.3 Create Service Utilization Feature
# -----
# Combine outpatient, emergency, and inpatient visits into a single aggregated
# ↪feature.

df['service_utilization'] = (
    df['number_outpatient'] +
    df['number_emergency'] +
    df['number_inpatient']
)
```

We created `age_mid` (age interval midpoints) and `service_utilization` (total prior visits) as simple feature transformations for modeling.

```
[32]: # 4.4 Save Processed Dataset
# -----
# Save the cleaned and feature-engineered dataset for use in modeling.

output_path = "../data/processed/diabetic_data_clean.csv"
df.to_csv(output_path, index=False)
print(f"Saved processed dataset to {output_path}")
```

Saved processed dataset to ../data/processed/diabetic_data_clean.csv

The processed dataset is saved for subsequent modeling and explainability analysis.


```
[33]: # 4.5 Basic Visualizations for Dataset Description
import matplotlib.pyplot as plt
import seaborn as sns

numeric_vars = ['time_in_hospital',
                 'num_lab_procedures',
                 'num_medications',
                 'service_utilization']

plt.figure(figsize=(12, 10))

for i, col in enumerate(numeric_vars, 1):
    plt.subplot(2, 2, i)
    sns.histplot(df[col], kde=True, bins=30)
    plt.title(f"Distribution of {col}")
    plt.xlabel(col)

plt.tight_layout()
plt.savefig(FIG_DIR / "fig2_distribution.png", dpi=300, bbox_inches="tight")
plt.show()
```



```
[34]: # -----
# SETUP: Import Libraries and Load Data
# -----
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

# Set a professional plot style
sns.set_style("whitegrid")
plt.rcParams['figure.dpi'] = 150 # High resolution for export

# NOTE: Adjust this path if your diabetes_cleaned.csv is in a different location
try:
    df = pd.read_csv("diabetes_cleaned.csv")
    print("Cleaned data loaded successfully.")
except FileNotFoundError:
    print("Error: 'diabetes_cleaned.csv' not found. Please check the file path.
    ↵")
    exit()

# -----
# 1. DISTRIBUTION OF KEY NUMERIC VARIABLES
# Supporting Section 3.4: Shows skewness and feature concentration
# -----
print("\nGenerating Figure 1: Distribution of Key Numeric Variables...")

key_numerics = ['time_in_hospital', 'num_lab_procedures', 'num_medications', ↵
    ↵'service_utilization']
titles = [
    'Time in Hospital (Days)',
    'Number of Lab Procedures',
    'Number of Medications',
    'Composite Service Utilization'
]

fig, axes = plt.subplots(2, 2, figsize=(12, 8))
axes = axes.flatten()

for i, col in enumerate(key_numerics):
    sns.histplot(df[col], kde=True, ax=axes[i], bins=30, color="#1f77b4")
    axes[i].set_title(f'Distribution of {titles[i]}', fontsize=12)
    axes[i].set_xlabel(titles[i])
    axes[i].set_ylabel('Count (Log Scale)')
```

```

        axes[i].set_yscale('log') # Use log scale for clearer visualization of
        ↪skewed data

plt.tight_layout()
plt.suptitle("Figure 3.1: Distribution of Key Numeric Variables in Cleaned"
    ↪Dataset", y=1.02, fontsize=16)
plt.savefig('figure_3_1_distributions.png')
plt.show()

# -----
# 2. READMISSION RATE BY HOSPITAL STAY DURATION
# Supporting Section 3.4: Confirms longer stays = higher risk (critical finding)
# -----
print("\nGenerating Figure 2: Readmission Rate by Hospital Stay Duration...")

# Calculate the mean readmission rate (1=readmitted) for each time_in_hospital
# group
readmission_rate_by_stay = df.groupby('time_in_hospital')['readmitted_binary'].mean().reset_index()

plt.figure(figsize=(8, 5))
sns.lineplot(
    data=readmission_rate_by_stay,
    x='time_in_hospital',
    y='readmitted_binary',
    marker='o',
    color='#d62728',
    linewidth=2
)

# Add a horizontal line to indicate the overall mean rate (11.16%)
overall_rate = df['readmitted_binary'].mean()
plt.axhline(overall_rate, color='grey', linestyle='--', label=f'Overall Rate'
    ↪({overall_rate*100:.2f}%)')

plt.title("Figure 3.2: Readmission Rate (%) by Hospital Stay Duration (Days)", ↪
    ↪fontsize=14)
plt.xlabel("Time in Hospital (Days)", fontsize=10)
plt.ylabel("30-Day Readmission Rate", fontsize=10)
plt.yticks([0.05, 0.10, 0.15, 0.20]) # Set fixed y-ticks for clarity
plt.ylim(0.05, 0.20)
plt.legend()
plt.savefig('figure_3_2_readmission_by_stay.png')
plt.show()

```

```

# -----
# 3. AVERAGE MEDICATIONS & TIME IN HOSPITAL BY AGE GROUP
# Supporting Section 3.4: Shows complexity/severity increases with age
# -----
print("\nGenerating Figure 3: Average Metrics by Age Group...")

# The 'age' column is already categorical (e.g., [50-60]), suitable for grouping
age_trends = df.groupby('age').agg(
    avg_medications=('num_medications', 'mean'),
    avg_time_in_hospital=('time_in_hospital', 'mean')
).reset_index()

# Explicitly define the correct chronological order for the age categories
age_order = ['[0-10)', '[10-20)', '[20-30)', '[30-40)', '[40-50)',
             '[50-60)', '[60-70)', '[70-80)', '[80-90)', '[90-100)']

# Ensure age groups are plotted in the correct order by setting the category order
age_trends['age'] = pd.Categorical(age_trends['age'], categories=age_order, ordered=True)
age_trends = age_trends.sort_values('age')

fig, ax1 = plt.subplots(figsize=(10, 6))

color1 = '#2ca02c' # Green for medications
sns.lineplot(data=age_trends, x='age', y='avg_medications', ax=ax1, marker='s', color=color1)
ax1.set_ylabel('Average Medications', color=color1)
ax1.tick_params(axis='y', labelcolor=color1)

# Create a second axis for Time in Hospital
ax2 = ax1.twinx()
color2 = '#ff7f0e' # Orange for time in hospital
sns.lineplot(data=age_trends, x='age', y='avg_time_in_hospital', ax=ax2, marker='o', color=color2)
ax2.set_ylabel('Average Time in Hospital (Days)', color=color2)
ax2.tick_params(axis='y', labelcolor=color2)

ax1.set_title("Figure 3.3: Average Clinical Metrics by Patient Age Group", fontsize=14)
ax1.set_xlabel("Age Group", fontsize=10)
plt.tight_layout()
plt.savefig('figure_3_3_age_trends.png')
plt.show()

```

```

# -----
# 4. CORRELATION HEATMAP (REFINED)
# Supporting Section 3.4 & 4.3: Justifies baseline model choice
# -----
print("\nGenerating Figure 4: Correlation Heatmap...")

# Select only numeric columns for correlation matrix
numeric_cols = df.select_dtypes(include=[np.number]).columns.tolist()
corr_df = df[numeric_cols].corr()

# 1. Focus only on correlations with the target variable, 'readmitted_binary'
target_corr = corr_df['readmitted_binary'].sort_values(ascending=False).
    ↪drop('readmitted_binary')

# 2. Select the top 10 absolute correlations for better visualization in the
# report
# We include 'service_utilization' even if it's not in the top 10 absolute (it
# ↪is in the top 5 here)
top_features = target_corr.abs().sort_values(ascending=False).index[:10].
    ↪tolist()
target_corr_top = target_corr.loc[top_features]


plt.figure(figsize=(5, 8))
sns.heatmap(
    target_corr_top.to_frame(),
    annot=True,
    fmt=".3f",
    cmap='vlag',
    cbar=False,
    linewidths=0.5,
    linecolor='black'
)
plt.title("Figure 3.4: Top 10 Feature Correlations with Readmission Target", ↪
    fontsize=14)
plt.yticks(rotation=0)
plt.tight_layout()
plt.savefig('figure_3_4_correlation_heatmap.png')
plt.show()

print("\nAll four required EDA figures generated and saved as PNG files.")
print("The files are: figure_3_1_distributions.png, ↪
    ↪figure_3_2_readmission_by_stay.png, figure_3_3_age_trends.png, ↪
    ↪figure_3_4_correlation_heatmap.png.")

```


Error: 'diabetes_cleaned.csv' not found. Please check the file path.

Generating Figure 1: Distribution of Key Numeric Variables...




Generating Figure 2: Readmission Rate by Hospital Stay Duration...

Figure 3.2: Readmission Rate (%) by Hospital Stay Duration (Days)




Generating Figure 3: Average Metrics by Age Group...

Figure 3.3: Average Clinical Metrics by Patient Age Group



Generating Figure 4: Correlation Heatmap...

Figure 3.4: Top 10 Feature Correlations with Readmission Target



All four required EDA figures generated and saved as PNG files.

The files are: figure_3_1_distributions.png, figure_3_2_readmission_by_stay.png, figure_3_3_age_trends.png, figure_3_4_correlation_heatmap.png.

1.6 5. Exploratory Data Analysis (EDA)

In this section, we visualize distributions, relationships, and trends in the cleaned dataset to gain insights that may inform feature importance and model design.

[35]: # 5.1 Dataset Overview

```
# -----
print(df.shape)
df.info()
df.describe(include='all')
```

(101763, 46)
<class 'pandas.core.frame.DataFrame'>
Index: 101763 entries, 0 to 101765
Data columns (total 46 columns):

#	Column	Non-Null Count	Dtype
0	race	101763 non-null	object
1	gender	101763 non-null	object
2	age	101763 non-null	object
3	admission_type_id	101763 non-null	int64
4	discharge_disposition_id	101763 non-null	int64
5	admission_source_id	101763 non-null	int64
6	time_in_hospital	101763 non-null	int64
7	num_lab_procedures	101763 non-null	int64
8	num_procedures	101763 non-null	int64
9	num_medications	101763 non-null	int64
10	number_outpatient	101763 non-null	int64
11	number_emergency	101763 non-null	int64
12	number_inpatient	101763 non-null	int64
13	diag_1	101763 non-null	object
14	diag_2	101763 non-null	object
15	diag_3	101763 non-null	object
16	number_diagnoses	101763 non-null	int64
17	metformin	101763 non-null	object
18	repaglinide	101763 non-null	object
19	nateglinide	101763 non-null	object
20	chlorpropamide	101763 non-null	object
21	glimepiride	101763 non-null	object
22	acetohexamide	101763 non-null	object
23	glipizide	101763 non-null	object
24	glyburide	101763 non-null	object
25	tolbutamide	101763 non-null	object
26	pioglitazone	101763 non-null	object
27	rosiglitazone	101763 non-null	object

```

28 acarbose           101763 non-null object
29 miglitol           101763 non-null object
30 troglitazone       101763 non-null object
31 tolazamide          101763 non-null object
32 examide             101763 non-null object
33 citoglipton         101763 non-null object
34 insulin              101763 non-null object
35 glyburide-metformin 101763 non-null object
36 glipizide-metformin 101763 non-null object
37 glimepiride-pioglitazone 101763 non-null object
38 metformin-rosiglitazone 101763 non-null object
39 metformin-pioglitazone 101763 non-null object
40 change               101763 non-null object
41 diabetesMed          101763 non-null object
42 readmitted           101763 non-null object
43 readmitted_binary    101763 non-null int8
44 age_mid              101763 non-null int64
45 service_utilization  101763 non-null int64
dtypes: int64(13), int8(1), object(32)
memory usage: 35.8+ MB

```

```

[35]:      race   gender      age admission_type_id dischargeDisposition_id
admission_source_id \
count      101763  101763    101763      101763.000000          101763.000000
101763.000000
unique        5      2      10                  NaN                  NaN
NaN
top      Caucasian Female [70-80)                  NaN                  NaN
NaN
freq      78370    54708    26066                  NaN                  NaN
NaN
mean      NaN      NaN      NaN      2.024017          3.715515
5.754459
std      NaN      NaN      NaN      1.445414          5.279919
4.064110
min      NaN      NaN      NaN      1.000000          1.000000
1.000000
25%      NaN      NaN      NaN      1.000000          1.000000
1.000000
50%      NaN      NaN      NaN      1.000000          1.000000
7.000000
75%      NaN      NaN      NaN      3.000000          4.000000
7.000000
max      NaN      NaN      NaN      8.000000          28.000000
25.000000

time_in_hospital  num_lab_procedures  num_procedures  num_medications

```

	number_outpatient	number_emergency	\			
count	101763.000000	101763.000000	101763.000000	101763.000000	101763.000000	101763.000000
101763.000000	101763.000000					
unique	NaN	NaN	NaN	NaN	NaN	NaN
NaN	NaN					
top	NaN	NaN	NaN	NaN	NaN	NaN
NaN	NaN					
freq	NaN	NaN	NaN	NaN	NaN	NaN
NaN	NaN					
mean	4.396018	43.095909	1.339691	16.021835		
0.369368	0.197842					
std	2.985092	19.674220	1.705792	8.127589		
1.267282	0.930485					
min	1.000000	1.000000	0.000000	1.000000		
0.000000	0.000000					
25%	2.000000	31.000000	0.000000	10.000000		
0.000000	0.000000					
50%	4.000000	44.000000	1.000000	15.000000		
0.000000	0.000000					
75%	6.000000	57.000000	2.000000	20.000000		
0.000000	0.000000					
max	14.000000	132.000000	6.000000	81.000000		
42.000000	76.000000					

	number_inpatient	diag_1	diag_2	diag_3	number_diagnoses	metformin	repaglinide	nateglinide	chlorpropamide	\
count	101763.000000	101763	101763	101763	101763.000000	101763				
101763	101763	101763								
unique	NaN	717	749	790	NaN	4				
4	4	4								
top	NaN	428	276	250	NaN	No				
No	No	No								
freq	NaN	6862	6752	11555	NaN	81776				
100224	101060	101677								
mean	0.635585	NaN	NaN	NaN	7.422649	NaN				
NaN	NaN	NaN								
std	1.262877	NaN	NaN	NaN	1.933578	NaN				
NaN	NaN	NaN								
min	0.000000	NaN	NaN	NaN	1.000000	NaN				
NaN	NaN	NaN								
25%	0.000000	NaN	NaN	NaN	6.000000	NaN				
NaN	NaN	NaN								
50%	0.000000	NaN	NaN	NaN	8.000000	NaN				
NaN	NaN	NaN								
75%	1.000000	NaN	NaN	NaN	9.000000	NaN				
NaN	NaN	NaN								
max	21.000000	NaN	NaN	NaN	16.000000	NaN				

	NaN	NaN	NaN	glimepiride acetohexamide glipizide glyburide tolbutamide pioglitazone rosiglitazone acarbose miglitol \			
count	101763	101763	101763	101763	101763	101763	101763
101763	101763	101763					
unique		4	2	4	4	2	4
4	4	4					
top		No	No	No	No	No	No
No	No	No					
freq	96572	101762	89078	91113	101740	94436	
95399	101455	101725					
mean		NaN	NaN	NaN	NaN	NaN	NaN
NaN	NaN	NaN					
std		NaN	NaN	NaN	NaN	NaN	NaN
NaN	NaN	NaN					
min		NaN	NaN	NaN	NaN	NaN	NaN
NaN	NaN	NaN					
25%		NaN	NaN	NaN	NaN	NaN	NaN
NaN	NaN	NaN					
50%		NaN	NaN	NaN	NaN	NaN	NaN
NaN	NaN	NaN					
75%		NaN	NaN	NaN	NaN	NaN	NaN
NaN	NaN	NaN					
max		NaN	NaN	NaN	NaN	NaN	NaN
NaN	NaN	NaN					

	troglitazone tolazamide examide citoglipton insulin glyburide-metformin glipizide-metformin \						
count	101763	101763	101763	101763	101763	101763	101763
101763							
unique	2	3	1	1	4		4
2							
top	No	No	No	No	No	No	No
No							
freq	101760	101724	101763	101763	47380		101057
101750							
mean	NaN	NaN	NaN	NaN	NaN	NaN	NaN
NaN							
std	NaN	NaN	NaN	NaN	NaN	NaN	NaN
NaN							
min	NaN	NaN	NaN	NaN	NaN	NaN	NaN
NaN							
25%	NaN	NaN	NaN	NaN	NaN	NaN	NaN
NaN							
50%	NaN	NaN	NaN	NaN	NaN	NaN	NaN
NaN							

75%	NaN	NaN	NaN	NaN	NaN	NaN
NaN						
max	NaN	NaN	NaN	NaN	NaN	NaN
NaN						
glimepiride-pioglitazone metformin-rosiglitazone metformin-pioglitazone						
change diabetesMed readmitted \						
count		101763		101763		101763
101763	101763	101763				101763
unique			2		2	
2	2	3				2
top		No		No		No
No	Yes	NO				
freq		101762		101761		101762
54754	78361	54861				
mean		NaN		NaN		NaN
NaN	NaN	NaN				
std		NaN		NaN		NaN
NaN	NaN	NaN				
min		NaN		NaN		NaN
NaN	NaN	NaN				
25%		NaN		NaN		NaN
NaN	NaN	NaN				
50%		NaN		NaN		NaN
NaN	NaN	NaN				
75%		NaN		NaN		NaN
NaN	NaN	NaN				
max		NaN		NaN		NaN
NaN	NaN	NaN				
readmitted_binary		age_mid	service_utilization			
count	101763.000000	101763.000000	101763.000000			
unique		NaN	NaN	NaN		
top		NaN	NaN	NaN		
freq		NaN	NaN	NaN		
mean	0.111602	65.966854	1.202795			
std	0.314878	15.941022	2.291805			
min	0.000000	5.000000	0.000000			
25%	0.000000	55.000000	0.000000			
50%	0.000000	65.000000	0.000000			
75%	0.000000	75.000000	2.000000			
max	1.000000	95.000000	80.000000			

```
[ ]: # 5.2 Class Imbalance Visualization
# -----
plt.figure(figsize=(5, 4))
sns.countplot(x='readmitted_binary', data=df)
```

```

plt.title("Distribution of 30-day Readmission Labels")
plt.xlabel("Readmitted within 30 days (1 = Yes, 0 = No)")
plt.ylabel("Count")
plt.tight_layout()
plt.savefig(FIG_DIR / "fig_readmission_distribution.png", dpi=300, u
    ↪bbox_inches="tight")
plt.show()

```

The dataset is imbalanced, with only ~11% of patients readmitted within 30 days. This imbalance highlights the need for resampling or class-weighting techniques in later modeling stages.

```

[ ]: # 5.3 Continuous Feature Distributions
# -----
plt.figure(figsize=(6,4))
sns.boxplot(x='readmitted_binary', y='time_in_hospital', data=df)
plt.title("Hospital Stay Duration vs Readmission")
plt.xlabel("Readmitted within 30 days (1 = Yes, 0 = No)")
plt.ylabel("Time in Hospital (days)")
plt.show()

```

```

[ ]: # 5.4 Feature Relationships (categorical features)
# -----
fig, axes = plt.subplots(1, 2, figsize=(12, 5))

sns.countplot(x='race', hue='readmitted_binary', data=df, ax=axes[0])
axes[0].set_title("Race vs Readmission")
axes[0].set_xlabel("Race")
axes[0].tick_params(axis='x', rotation=45)

sns.countplot(x='admission_type_id', hue='readmitted_binary', data=df, u
    ↪ax=axes[1])
axes[1].set_title("Admission Type vs Readmission")
axes[1].set_xlabel("Admission Type ID")
axes[1].tick_params(axis='x', rotation=45)

plt.tight_layout()
plt.show()

plt.figure(figsize=(6, 4))
sns.boxplot(x='readmitted_binary', y='service_utilization', data=df)
plt.title("Service Utilization vs Readmission")
plt.xlabel("Readmitted within 30 days (1 = Yes, 0 = No)")
plt.ylabel("Total Prior Visits")
plt.show()

```

Features A1Cresult and max_glu_serum were dropped earlier due to excessive missingness (>80%). We instead examine relationships between readmission and remaining features such as race, admission type, and service utilization.